

Day 1 Competition Checklist



Guideline source selected and legal usability confirmed



PDF parsed with structure preserved (tables, headings intact)



Chunking strategy chosen and applied consistently



Embedding model selected and justified



Vector index built with document/section/page metadata attached



A test query returns a sensible, relevant chunk

Common Day 1 Pitfalls to Avoid



Scope creep

Trying to ingest 4+ guideline documents instead of 1-2



Skipping the legal check

Parsing a source before confirming it's actually usable



Losing table structure

Treating dosage tables as flat paragraph text



No metadata plan

Indexing chunks without page/section tracking — impossible to cite later

What “Done” Looks Like at The End

- 1 A working, queryable vector index built from your chosen source
- 2 Every entry traceable to its exact document, section, and page
- 3 A short written note on why you chose your chunking strategy and embedding model
- 4 Confidence that tomorrow's retrieval work has clean, well-structured data to search